

Lab 3

Linear Regression

```
import numpy as np
import matplotlib.pyplot as plt
```

```
def estimate_coef(x, y):
    n = np.size(x)
    m_x = np.mean(x)
    m_y = np.mean(y)
    SS_xy = np.sum((x - m_x) * (y - m_y))
    SS_xx = np.sum((x - m_x) ** 2)
    b_1 = SS_xy / SS_xx
    b_0 = m_y - b_1 * m_x
    return (b_0, b_1)
```

```
def plot_regression_line(x, y, b):
    plt.scatter(x, y, color="m", marker="o", s=30)
    y_pred = b[0] + b[1] * x
    plt.plot(x, y_pred, color="g")
    plt.xlabel('x')
    plt.ylabel('y')
    plt.title("Linear Regression")
    plt.show()
```

```
file_path = input("Content of marketing.csv")
df = pd.read_csv(file_path)
```

```
x = df.iloc[:, 0].values
y = df.iloc[:, 1].values
```

```
b = estimate_coef(x, y)
print("Estimated coefficients: \nb_0 = %b[0]"
      "\nb_1 = %b[1]" % b)
plot_regression_line(x, y, b)
```


Output

$$b_0 = 7.03$$

$$b_1 = 0.04$$

Multiple Linear Regression

data = {

"Feature 1": [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],

"Feature 2": [2, 3, 5, 7, 11, 13, 17, 19, 23, 25],

"Feature 3": [3, 6, 9, 12, 15, 18, 21, 24, 27, 30],

"Target": [5, 9, 15, 22, 31, 41, 53, 66, 80, 96]

}

df = pd.DataFrame(data)

x = df.drop(columns=["Target"]).values

y = df["Target"].values.reshape(-1, 1)

x = np.hstack((np.ones((x.shape[0], 1)), x))

beta = np.linalg.solve((x.T @ x + 0.01 * np.identity(x.shape[1])), x.T @ y)

y_pred = x @ beta

mse = np.mean((y - y_pred)**2)

total_variance = np.sum((y - np.mean(y))**2)

explained_variance = np.sum((y - y_pred - np.mean(y))**2)

r2 = explained_variance / total_variance

print("Model coefficients:", beta[1:].flatten())

print("Intercept:", beta[0][0])

```
print("Mean Squared Error:", mse)
print("R-squared Score:", r2)
```

Output

Model coefficients: [0.04 3.31 0.12]
Intercept: -3.15

Mean Squared Error: 5.36

R-squared Error: 0.99.

Logistic Regression.

```
def sigmoid(z):
    return 1 / (1 + np.exp(-z))
```

```
def compute_cost(X, y, theta):
    m = len(y)
    h = sigmoid(X @ theta)
    cost = (-1/m) * np.sum(y * np.log(h) +
                             (1 - y) * np.log(1 - h))
    return cost
```

```
def gradient_descent(X, y, theta, alpha,
                     iterations):
    m = len(y)
    cost_history = []
```

```
    for i in range(iterations):
        gradient = (-1/m) * X.T @ (sigmoid
                                     (X @ theta) - y)
        theta = theta - alpha * gradient
        cost_history.append(compute_cost(X, y, theta))
```

```
    return theta, cost_history
```



```
def predict(x, theta):
    return sigmoid(x @ theta) > 0.5
astype(int)
```

```
np.random.seed(42)
X = np.random.rand(100, 1) * 10
y = (X > 5).astype(int).ravel()
```

```
X_b = np.c_[np.ones((X.shape[0], 1)), X]
```

```
theta = np.zeros(X_b.shape[1])
alpha = 0.1
iterations = 1000
```

```
theta, cost_history = gradient_descent(X_b,
                                         y, theta, alpha, iterations)
```

```
y_pred = predict(X_b, theta)
```

Output

Accuracy: 0.97

[0, 0, 1, 0, 0, 1, 1]

Subal B
24/3/25